

Heuristic Analysis

Problem 1

This problem is by far the ‘easiest’ to solve (from computational point of view). The search space consists of 2^{12} states. The following table presents the performance data:

Search Method	Expansions	Goal Tests	New Nodes	Plan Size	Time
breadth_first_search	43	56	180	6	0.031
breadth_first_tree_search	1458	1459	5960	6	0.993
depth_first_graph_search	21	22	84	20	0.014
depth_limited_search	101	271	414	50	0.095
uniform_cost_search	55	57	224	6	0.384
recursive_best_first_search	4229	4230	17023	6	2.960
greedy_best_first_graph_search	7	9	28	6	0.005
A* with h1	55	57	224	6	0.039
A* with h_ignore_preconditions	41	43	170	6	0.044
A* with h_pg_levelsum	11	13	50	6	2.234

One optimal plan consists of the following (6) actions:

1. Load(C1, P1, SFO)
2. Load(C2, P2, JFK)
3. Fly(P1, SFO, JFK)
4. Fly(P2, JFK, SFO)
5. Unload(C1, P1, JFK)
6. Unload(C2, P2, SFO)

The best performer is **greedy_best_first_graph_search**. It found optimal plan using far fewer operations. But was it luck?

Problem 2

The search space for this problem is significantly larger, yet still manageable. It consists of 2^{27} states. One key difference from the previous problem is that I wasn’t able to produce results for all searches, due to ‘waiting too much’. Those that are available are presented in the following table:

Search Method	Expansions	Goal Tests	New Nodes	Plan Size	Time
breadth_first_search	3343	4609	30509	9	11.708
depth_first_graph_search	624	625	5602	619	3.061
uniform_cost_search	4852	4854	44030	9	38.310
greedy_best_first_graph_search	990	992	8910	21	6.166

Search Method	Expansions	Goal Tests	New Nodes	Plan Size	Time
A* with h1	4852	4854	44030	9	39.245
A* with h_ignore_preconditions	1506	1508	13820	9	12.527
A* with h_pg_levelsum	86	88	841	9	258.185

One optimal plan (9 actions) is the following:

1. Load(C1, P1, SFO)
2. Load(C2, P2, JFK)
3. Load(C3, P3, ATL)
4. Fly(P2, JFK, SFO)
5. Unload(C2, P2, SFO)
6. Fly(P1, SFO, JFK)
7. Unload(C1, P1, JFK)
8. Fly(P3, ATL, SFO)
9. Unload(C3, P3, SFO)

Maybe it was luck? The performance from `greedy_best_first_graph_search` was great but the plan that was found wasn't optimal. Our winner this time is `breadth_first_search`. Note that it did pretty well on **Problem 1** as well.

Problem 3

The state space now grows to 2^{32} states. The time required to run a search is now much larger than before as shown by the following results:

Search Method	Expansions	Goal Tests	New Nodes	Plan Size	Time
breadth_first_search	14663	18098	129631	12	89.299
depth_first_graph_search	408	409	3364	392	1.476
uniform_cost_search	18235	18237	159716	12	329.735
greedy_best_first_graph_search	5614	5616	49429	22	87.676
A* with h1	18235	18237	159716	12	327.593
A* with h_ignore_preconditions	5118	5120	45650	12	74.635
A* with h_pg_levelsum	408	410	3758	12	1713.060

An optimal plan with 12 actions:

1. Load(C2, P2, JFK)
2. Fly(P2, JFK, ORD)
3. Load(C4, P2, ORD)
4. Fly(P2, ORD, SFO)
5. Unload(C4, P2, SFO)
6. Load(C1, P1, SFO)
7. Fly(P1, SFO, ATL)

8. Load(C3, P1, ATL)
9. Fly(P1, ATL, JFK)
10. Unload(C3, P1, JFK)
11. Unload(C1, P1, JFK)
12. Unload(C2, P2, SFO)

Finally, we have a winner that uses a heuristic. In this case it is **A* search with h_ignore_preconditions**. Note that **breadth_first_search** is still performing quite well and obtains an optimal result.

Conclusion

The analysis of the different problems did not provide a clear winner that is best in all situations. But why using a heuristic isn't always better? As long as the search space is small(ish), as in **Problem 1 and 2**, visiting every state is cheap and non heuristic methods are effective. Additionally, they are easier to understand/implement. When that space becomes large, though, this technique proves to be quite inefficient. Now, computing a heuristic has smaller cost compared to visiting every node in the planning graph.